

Linux Notes

CONTENTS

I	Introduction	1
II	Bash	3
III	Environment & Configuration	4
III-A	Environment variables	4
III-B	Configuration files	5
IV	System Administration	6
IV-A	Package Management	7
IV-B	Devices	8
IV-C	C & Python	9
IV-D	Accelerated computing	10

I. INTRODUCTION

- GNU:
 - GNU is a Unix-like operating system
 - * it is a collection of many programs: applications, libraries, developer tools, even games
 - * GNU is free software that respects users' freedom
 - GNU stands for GNU's Not Unix
 - initiated by *Richard Stallman*
 - * also initiated Free Software movement & Emacs
 - GNU C, is referred to as GCC
 - GNU General Public License (GPL)
- Linux:
 - created in 1991, by Linus Torvalds
 - literally, *Linux* is the name of the OS's kernel
 - in popular usage, Linux includes the whole ecosystem
 - * including accompanying GNU programs
 - * other programs
- Three levels of Linux system:
 - 1) hardware
 - 2) kernel
 - 3) user processes
- Kernel:
 - the *kernel* is the core of OS
 - the kernel is software in memory that tells the CPU what to do
 - the kernel manages HW and is the interface between HW and running program
 - the kernel manages the following:
 - * system calls from processes, see next
 - * managing processes
 - * memory management
 - * device drivers
- User processes:
 - *user processes* are the running programs that the kernel manages
 - these processes collectively form the *user space*
- Kernel versus user modes/spaces:
 - the kernel runs in *kernel mode* & has unrestricted access
 - the user processes run in *user modes* which are restricted
 - the part of the main memory that user processors can access is called *user space*
 - the area that only the kernel can access is called the *kernel space*
- Ubuntu:
 - *Ubuntu* is a Linux distribution
 - code name is *xenial*
- User & root user:
 - a *user* is an entity that runs processes & own files
 - * users exist to support permissions and boundaries
 - a user is associated with a *username*
 - the kernel identifies users by *userids*
 - every user process has a *user owner*
 - the *root user* is an exception because they can interfere with other users,
 - * also known as *superuser*
- Filesystem:
 - Linux has a single filesystem tree
 - * this is unlike Windows
 - * storage devices are mounted at various points on the tree
 - filenames that begin with . symbol are hidden
 - filenames and commands are case sensitive
 - no concept of file extension
 - do not imbed spaces in filenames
- Permissions:
 - a file's *mode* summarizes the files permission
 - the mode has the following format


```
trwxrwxrwx
```
 - t stands for *type*
 - * 't=-' means file
 - * 't=d' means directory
 - * 't=b' means *block*
 - * 't=c' means *character*
 - * 't=p' means *pipe*
 - * 't=s' means *socket*
 - * 'b', 'c', 'p', or 's' imply a device
 - the 9 remaining fields are partitioned into 3 sets:
 - * permissions for user (u)

- * permission for group (p)
- * permission for other (o)
- each set has three fields
 - * read (r)
 - * write (w) and
 - * executable (x)
- absence of a permission is denoted by '-'
- directories need to be executable to be able to be accessible
- Running an executable file:
 - make sure file permission 'x' is turned on
 - `ttype ./fn`
- Common directories:
 - `/bin` contains binaries, or executables, that must be present for the system to boot and run
 - `/boot` contains Linux Kernel, initial RAM disk image (for drivers needed at boot time) and the boot loader
 - `/dev` contains device nodes
 - `/etc`, contains system configuration files
 - * shell scripts, password, boot, device and networking setup files
 - `/home`, contain regular user directories
 - `/lib`, abbreviation for library, DLL equivalent
 - `/media`, mount points for removable media such as USB drives and CD-ROMs
 - `/mnt` is a generic mount point under which you mount your file-systems or devices
 - `sys` device and system interface
 - `sbin`, system bin, system executables, for system administrators
 - `/usr` contains bulk of Linux system programs & support files used by users
 - * many of the directory names are the same as the root directory, and hold same type of data
 - * `/usr/include` holds header files by the C compiler
 - * `/usr/info` contains GNU info manuals
 - * `/usr/local` is where administrators store their own software
 - * `/usr/man` contains manual pages
 - `/var`, variable directory, where runtime information is stored
- Text editors:
 - `emacs`
 - `nano` #basic editor
 - `vi`
 - `less`
 - * see the contents of a file one screen at a time
 - * press spacebar for next screen
 - * press b for previous screen
 - * press /text to search for text forward
 - * press q to exit
 - * can redirect to less command | less
- Remote log in:
 - OpenSSH is a freely available version of the *Secure*

Shell (SSH) protocol family of tools for remotely controlling, or transferring files between, computers

- X windows system:
 - GUI (Graphic User Interface) uses a terminal emulator to interact with the shell
 - WIMP stands for windows, icons, menus, pointer
 - a *windowing system* (or window system) is a type of GUI which implements the WIMP paradigm for a user interface
 - * can think of a windowing system as a device driver
 - the *X Window System* (X11, or shortened to simply X) is a windowing system for bitmap displays, common on UNIX-like computer operating systems
 - think of X as sort of the kernel of the desktop that manages anything from generating windows to configuring displays to handling input devices
 - * for example, X windows allows copying using mouse right and middle buttons
 - there is an *X server* and an *X client*
- Desktop environment: Gnome:
 - a *desktop environment* (DE), is a collection of software designed to give functionality and a certain look and feel to an operating system
 - * important parts of a DE are the window manager (WM) and the file manager
 - * DE provides utilities to set wallpapers and screen-savers, display icons on the desktop, and perform some administrative tasks
 - * in Ubuntu all DE's uses X windows system
 - Gnome is a desktop environment and is part of GNU Project
 - * Gnome3 is the current version
 - Ubuntu tracks Gnome, but uses *Unity shell* rather than Gnome shell

II. BASH

- Shell:
 - a shell is a program that takes keyword commands and passes it to the OS
 - a shell is a *command line interface* (CLI)
 - general command structure is
`command -options arguments`
 - options could be
 - * a single character (e.g. -l) or
 - * long option, (e.g. --reverse)
 - a *shell script* are text files that contain a sequence of shell commands
- Shell prompt:
 - commands are entered at the *shell prompt*
 - format is usually `user@machinename` followed by current working directory
 - \$ ending implies a regular user, # implies a superuser
 - a terminal is initialized to its home directory
- Bourne shell:
 - there are many different Unix shells
 - all shells derive features from Bourne shell
 - Bourne shell program is located at `/bin/sh`
 - Bourne shell was developed at Bell Labs
- bash:
 - the shell used in Linux/Ubuntu is *bash*
 - bash stands for Bourne-Again Shell
 - bash is an enhanced version of Bourne shell
 - `/bin/sh` is a link to `bash /bin/bash`
- Streams & standard I/O:
 - processes use I/O *streams* to read and write data
 - each process may be associated with
 - * *stdin* or *standard input*
 - * *stdout* or *standard output*
 - * *standard error*
- Redirection & pipe:
 - use the redirection character to send the output of a command to a file
`command > fn` #fn overwritten
`command >> fn` #appended to fn
 - to send stdout of a command to the stdin another command
`command1 | command2`
- Control commands:
 - ^ + Alt + T starts a terminal
 - ^ + D stops a current standard input
 - ^ + C terminated a program
- history:
 - up arrow shows previous command
 - `history` command lists history
- Filesystem navigation symbols:
 - the `.` symbol refers to the current working directory
 - * when `./` is pre-pended to a filename, then it is searched only in current directory
 - the `..` symbol refers to the parent working directory
 - the `~` symbol refers to home directory
- Wildcards:
 - wildcards can be used with any command that that accepts fn as argument
 - * symbol represents any characters
 - ? symbol represents a single character
 - [characters] represents any character ∈ set
 - * e.g. [abc]
 - [[: class]] represents any character that ∈ class
 - * e.g. [[: alnum]], [[: alpha]], [[: digit]], [[: lower]], [[: upper]]
- help with commands:
 - the *manual* command is `man`
 - * the manual has eight sections

```
command-name --help #basic help
man command      #more detailed
man -k keyword   #to search
info command     #more complex
```
- Simple commands:


```
date
cal      # calender
clear    # brings cursor to top of screen
exit     # exits terminal
```
- Filesystem navigation commands:


```
pwd      # print working directory
ls       # list
ls -a    # all files
ls -l    # long format
ls -lt   # sorted- last modified time
cd       # change directory
cd       # to home directory
```
- File & directory manipulation:


```
cp       # copy
cp item1 item2 # 1->2
cp item1 item2 dir #1,2 -> d
cp -R ...    # recursive, with dir
mv        # move
mkdir     # make directory
rmdir     # remove directory
rm -rf *   #recursively remove files/dir
rm        # remove
touch fn   # creates fn
file fn    # tells type of fn
diff fn1 fn2
```
- echo text
 - echo prints its argument to stdout
 - `echo $X` prints value of X variable
- cat fn
 - cat reads files sequentially, writing them to stdout
 - the name cat is derived from its function to concatenate files

- `grep text dir_fn`
 - looks for a text in directory or filename `dir_fn`
 - `-i` option makes text case-insensitive
- `find dir -name fn -print`
 - to find `fn` in directory `dir`
- `passwd`
 - to change password
- Processes:
 - a process can be executed in background by ending the command with `&`
 - each process on the system has a numeric *process ID*, or *PID*
 - `ps` reports a snapshot of the current processes


```
ps -elf #e=every, f=full-format
ps x    #all
ps -elf | grep lightdm
kill pid #terminates, last resort
```
- Changing file permission: `chmod`

```
chmod g+r fn #read permission to group
chmod o-w fn #removes write permission
chmod 644 fn #octal, absolute change
```
- Link & alias: `ln` & `alias`
 - a link can be hard (`ln`), or symbolic (`ln -s`)
 - a *symbolic link* is a file that points to another file or directory effectively creating an alias
 - symbolic links offer quick access to obscure directory paths
 - to create a symbolic link, type


```
ln -s target link-name
```

where the `target` is where the link is pointing to,

 - `alias` lists the current aliases
- Compress and archive:


```
gzip fn #compress, 1 fn, --> fn.gz
gunzip fn.gz #fn.gz --> fn
tar cvf archive.tar fn1 fn2 #fn->ar
tar xvf archive.tar #archive->fn
tar tvf archive.tar #table-of-content
```

 - `.tar` suffix is convention not requirement
 - `c` flag means create
 - `x` flag means extract
 - `f` flag should precede archive file
 - `p` flag preserves permissions
 - an archive could be compressed `fn.tar.gz`
- `zip` & `unzip`:
 - to `zip` all files in a directory type


```
zip fn *
```
 - `unzip fn`
 - `fn` does not include extension
- `ldd`
 - `ldd` invokes the standard dynamic linker
 - the library `GLIBC` is the standard dynamic linker

III. ENVIRONMENT & CONFIGURATION

- Shell variable:
 - the shell can store temporary variables, called *shell variables*,


```
Y=10    #to set, no space
Y="$Z"  #assignment
echo $Y
```
 - Environment variable & `export`
 - an *environment variable* is similar to a shell variable but it is not specific to a shell,
 - the OS passes the environment variables to programs that the shell runs,
 - assign an environment variable using the `export` command,
 - exporting a variable makes the variable available to all sub-shells and processes created by that shell,
 - it does not make it available everywhere in the system, only by processes created from that shell.
 - `set`
 - to set or unset options and positional parameters,
 - without options displays both the shell and environment variables.
 - `printenv`
 - displays only environment variables,
 - `printenv Y` lists the value of `Y`.
- #### A. Environment variables
- `PATH`
 - a *command path* is a list of system directories that the shell searches when trying to locate a command,
 - `PATH` is a special environment variable that contains the command path,
 - can add a directory `dir` first or last as follows


```
PATH=dir:$PATH
PATH=$PATH:dir
```
 - `LD_LIBRARY_PATH`
 - dynamic link library path.
 - `DISPLAY`
 - X window system variable,
 - is the name of the display if running a graphical environment,
 - general format is


```
DISPLAY = hostname:D.S
```

 - * omitted hostname means local host,
 - * `D` means display number,
 - * `S` means screen number.
 - for example, `:0` means its the first display generated by X server,
 - `TERM`
 - name of terminal type
 - `xterm`.
 - `HOME`
 - pathname of home directory,

- /home/ara.
- USER
 - username,
 - ara.
- PS1
 - stands for *prompt string* 1,
 - built-in shell variable,
 - defines the contents of the shell prompt,
 - initially defined in /etc/profile,
 - then modified in /home/ara/.bashrc.
- SHELL
 - name of shell program being used,
 - /bin/bash.

- Activating changes: `source`
 - changes to startup files will not take effect until the terminal session is closed and start a new one,
 - `source` is a bash shell built-in command that executes the content of the file passed as argument,


```
source fn
```
 - `source` has a synonym in the symbol `'.'` (dot)


```
. fn
```

B. Configuration files

- Startup files:
 - when a user logs on, the bash program initiates a series of configuration scripts called *startup files*,
 - there are global scripts followed by user scripts,
 - * global startup files in /etc directory, define the default environment shared by all users,
 - * user startup files, /home/ara, can be used to extend or override settings in global configuration script,
 - in addition, there are two kinds of shell sessions: login shell sessions & non-login shell sessions,
 - * a *login session* prompts for user name & password,
 - * a *non-login session* typically occurs when a new terminal is launched.
- /etc/profile
 - global startup file for login sessions,
 - in interactive session it launches


```
/etc/bashrc,
```
 - also execute the scripts


```
/etc/profile.d/*.sh.
```
- /home/ara/.profile
 - user startup file for login sessions,
 - in interactive session it launches


```
/home/ara/.bashrc,
```
- /etc/bash.bashrc
 - global startup file for non-login sessions,
 - is used to set environmental items for a users shell,
 - is executed for both interactive and non-interactive shells,
- /home/ara/.bashrc
 - user startup file for non-login sessions,
 - I have CUDA initialization, and my aliases.
- /etc/X11/xorg.conf
 - the X configuration file provides a means to configure the X server,
 - the NVIDIA driver includes a utility called `nvidia-xconfig`, which is designed to make editing the X configuration file easy,
 - for diagnosis read /var/log/Xorg.0.log
 - * lines that start with (EE) indicate error.

IV. SYSTEM ADMINISTRATION

- Version I am using:
 - Ubuntu version 16.04.1 LTS
 - kernel version 4.4.0 – 47–generic
 - GCC 5.4.0
 - GLIBC 2.23
 - Python 2.7.12
- `lsb_release -a`
 - `lsb` means Linux Standard Base
 - tells you about OS details
- `uname:`
 - `uname -m` #32 vs 64 bit
 - `uname -r` #kernel version
- `sudo`
 - super user, administrative rights
 - precede any command with `sudo`
 - can enter super user mode by typing `sudo su`
- Storage management:
 - `df -h` # disk free, h=human
 - `du -h` # disk (dir) usage
 - `free` # free memory
- Dependency clashing:
 - if a package requires a shared resource (shared library), it is said to have a *dependency*
 - dependencies can be problematic when two different applications require incompatible versions of the same dependency
- Dependency environments:
 - to address dependency clashing, some package distribution managers have software that create environments, inside of which specific versions of software can be maintained independent of those contained in other environments
- Operating-system-level virtualization:
 - *OS-level virtualization* is a server virtualization method in which the kernel of an OS allows the existence of multiple isolated user-space instances, instead of just one
 - such instances are sometimes called *containers*, *software containers*, *virtualization engines* or *jails*
 - containers may look and feel like a real server from the point of view of its owners and users
 - *Docker* is an open-source project that automates the deployment of applications inside software containers
 - recommended to use TensorFlow in a container if running over multiple computers
- `dmesg`
 - `dmesg` stands for *display message*
 - `dmesg` is used to examine or control the kernel ring buffer
 - * `dmesg` lists the message buffer of the kernel
- the output typically contains the messages produced by the device drivers
- when the computer fails unexpectedly after reboot can use this command
- Boot
 - to change boot parameters check `etc/default/grub`

A. Package Management

- Package management & distributors:
 - *package management* is a method of installing and maintaining software on the system
 - one of the most important factors of a distribution is the quality of the packaging system
 - most distributions fall into one of two camps of packaging technologies:
 - * *Debian* style (.deb) that includes Ubuntu
 - * Red Hat style (.rpm)
- Package file:
 - the basic unit of software in a packaging system is the *package file*
 - a package file is a compressed collection of files that comprise the software package
 - the package file may include programs, data files, metadata, pre- and post-installation scripts, etc.
- Repositories:
 - most packages today are created by the distribution vendors
 - a distribution makes packages available to the user in central *repositories*
 - a distribution may also have related third-party repositories
 - package management system provide some method of dependency resolution
- Ubuntu package manager:
 - dpkg is a low level manager
 - apt, or *advanced package tool*, is high-level manager
 - * also stands for *aptitude*?
- dpkg


```
dpkg -l #lists installed packages
dpkg -i package #install
dpkg -r package #remove
dpkg --status package, #installed?
dpkg --search fn #package for fn
```
- apt


```
apt-get update #updates list only
apt-get upgrade #does installation
apt-get install xyz
apt-get remove package
apt autoremove #removes unnecessary
apt-cache search search-string
apt-cache show package
apt list --installed #installed packages
```
- PPA
 - a *personal package archive*, or PPA, is a software repository for uploading source packages to be built and published as an apt repository
 - PPA offers stable proprietary Nvidia graphics driver updates, without updating other libraries to unstable versions
 - to add graphics driver PPA:

```
sudo add-apt-repository
ppa:graphics-drivers/ppa
```

B. Devices

- Device drivers:
 - device drivers are part of the Linux kernel
 - * not in repository
 - * usually in /dev directory
 - * base path for devices is /sys/devices
 - the kernel presents many of the I/O device interfaces to user processes as files
 - exceptions when the driver is not in kernel:
 - * device is too new
 - * device is exotic, can compile
 - * HW vendor is hiding something
- `ubuntu-drivers devices`
 - Ubuntu specific command that lists devices
- `lshw`
 - `lshw` stands for *list hardware*
 - `lshw` extracts information about the hardware configuration of the machine


```
lshw -class display
lshw -short      #lists H
```
- `lspci`
 - `lspci` displays information about PCI buses in the system, and about the devices connected to them


```
lspci -v | grep -i nvidia
sudo update-pciids #update PCIe dev
```
- Display manager: LightDM
 - a *display manager* is responsible for starting the display server and loading the desktop after you type in your username and password
 - * display manager is not the same thing as a window manager or a display server
 - *LightDM* stands for *lightweight display manager*
 - LightDM is the default display manager for Ubuntu
 - * LightDM is an X display manager
 - * it starts the X servers, user sessions and greeter
 - the login window is known as the *greeter*
 - * LightDM offers separate greeter packages
 - * the default greeter in Ubuntu is *Unity Greeter*
 - display manager `lightdm` should be temporarily turned off before making changes to this config file and then restored
- `service`
 - can use `service stop` and `start` services temporarily


```
service --status-all
service lightdm status
sudo service lightdm stop
sudo service lightdm start
sudo service lightdm restart
```
- `update-alternatives`
 - creates, removes, maintains and displays information about the symbolic links,
 - `--config fn` shows available alternatives for a

link group and allow the user to interactively select which one to use,

- monitor showed the right resolution after I entered



```
sudo update-alternatives --config
x86_64-linux-gnu-gl_conf #debug
```


C. C & Python

- GCC
 - GNU compiler for C++,
 - to find version type `gcc -v`.
- GLIBC
 - the GNU C Library is the GNU Project's implementation of the C standard library,
 - despite its name, it now also directly supports C++,
 - to find version type `ldd --version`.
- Python:
 - start a Python session by typing `python`
 - to find version type `python --version`
 - in `/usr/local/lib/python2.7`
- python-pip
 - `pip` stands for *Installs Python* (recursive)
 - `pip` is a package management system to install and manage software packages written in Python
 - `apt-get install python-pip`
 - `python -m pip install --upgrade pip`
 - `pip install --upgrade pip`
 - `pip install --user scipy`
 - `pip install --user matplotlib`
 - `pip install --user sympy`
 - can use `pip` to install basic TensorFlow
- python-dev
 - `python-dev` is the package that contains the header files for the Python C API
 - `python-dev` is used by `lxml` because it includes Python C extensions for high performance
 - `apt-get install python-dev`
- python-numpy
 - `apt-get install python-numpy`
 - `pip install --user numpy`
- python-wheel
 - wheels are the new standard of python distribution
 - support is offered in `pip ≥ 1.4`
 - a wheel is a ZIP-format archive with a specially formatted filename and the `.whl` extension
 - wheel is designed to contain all the files that is very close to the on-disk format
 - `apt-get install python-wheel`
- Anaconda:
 - *Anaconda* is an open source distribution of the Python and R programming languages for large-scale data processing, predictive analytics, and scientific computing, that aims to simplify package management and deployment
 - <https://www.continuum.io/why-anaconda>
 - I have used Anaconda as a Python compiler
 - both Python versions are included in Anaconda
 - *Conda* is Anaconda's package manager application that quickly installs, runs, and updates packages and their dependencies
 - * *Conda* is also an environment manager application
- PyCharm
 - *PyCharm* is an Integrated Development Environment (IDE) used for programming in Python,
 - <https://www.jetbrains.com/pycharm/>
 - in Windows, used Anaconda compiler in PyCharm.
- Jupyter Notebook:
 - formerly *iPython Notebook*
 - interactively write documents that include code, text, output and *LateX*
- Python environments to address dependency clashing:
 - recommended to use a Python dependency environment when running TensorFlow on a single computer
 - two possibilities include:
 - * for the standard distributions use *Virtualenv*
 - * Anaconda comes with a built-in environment system
- python-virtualenv
 - to create a virtual Python environment
 - `apt-get install python-virtualenv`
 - create a directory, e.g. `env`, that contains this environment
 - create environment using `virtualenv` command
 - `virtualenv --system-site-packages ~/env/tensorflow`
 - activate environment using `source` command
 - `source ~/env/tensorflow/bin/activate`
 - environment should be active when installing with `pip`
 - exit environment by the command `deactivate`

D. Accelerated computing

- Installation hierarchy:



(1)

- NVIDIA GPU cards & drivers:
 - Titan X (Pascal), GeForce 10 series, GP102
 - GeForce GT 710B, GK208
 - installed in `/usr/lib/nvidia-367`
 - `nvidia-smi` shows GPUs and driver
 - driver version can be found as


```
cat /proc/driver/nvidia/version
```
 - `nvidia-xconfig` configures `xorg.conf`
- CUDA installation:
 - download from <https://developer.nvidia.com/cuda-downloads>
 - * use `deb(local)` option which is a large package
 - to install downloaded package


```
sudo dpkg -i fn.deb
sudo apt-get update
sudo apt-get install cuda
```
 - `libcupti-dev`
 - * the `libcupti-dev` library is the NVIDIA CUDA Profile Tools Interface
 - * this library provides advanced profiling support
 - * to install type


```
sudo apt-get install libcupti-dev
```
 - `nvcc`
 - * the `nvcc` command runs the compiler driver that compiles CUDA programs
 - * `nvcc` calls the GCC compiler for C code, & the NVIDIA PTX compiler for the CUDA code
 - * `nvcc --version` prints CUDA version

- CUDA

- installed in `/usr/local/cuda`
 - * `CUDA_HOME=/usr/local/cuda-8.0`
- `CUDA_VISIBLE_DEVICES`
 - * NVIDIA environment variable that specifies PCI device 0 (Titan) to be a CUDA device
 - * defined it in `/home/ara/.bashrc`
- `cudaSetDevice(a)` sets GPU *a* for CUDA computing
- `cuda_devices` is an alias that shows active CUDA devices
- after installation, query CUDA device by typing


```
/usr/local/cuda-8.0/samples/bin/...
x86_64/linux/release/deviceQuery
```

- cuDNN

- separate add-on designed for DNN
- download from <https://developer.nvidia.com/cudnn>
- inside `/usr/local/cuda` directory type


```
tar xvzf fn.tgz
cp cuda/include/cudnn.h include
chmod a+r include/cudnn.h
cp cuda/lib64/libcudnn* lib64
chmod a+r lib64/libcudnn*
```

- tensorflow

- TensorFlow is open source code provided by Google
- if using GPU, then build TensorFlow from source
- from home directory, clone latest TensorFlow source code from GitHub

```
git clone --recursive-submodules
https://github.com/
tensorflow/tensorflow
cd tensorflow
./configure
```

- Bazel

- * TensorFlow uses *Bazel* to build executable code from source code
- * Bazel is open source code provided by Google
 - Bazel is similar to *make*
- * install Bazel
- * create TensorFlow executable


```
bazel build -c opt --config=cuda
//tensorflow/tools/pip_package:
build_pip_package
```

- Python packages TensorFlow

- * dependency clashing prevention is done through Python